

Stats 598z: Homework 6

Due before class Thursday, Apr 8

Important:

R code, tables and figures should be part of a single .pdf or .html files from R Markdown and knitr. See the class reading lists for a short tutorial.

Include R commands for all output unless explicitly told not to.

If you collaborated with anyone else, mention their names and the nature of the collaboration

1 Problem 1: LASSO [40pts]

- (a) Write a function `gen_data` to generate a training dataset (X, Y) . Your function should take in 4 arguments, `n`, `p`, `sparsity` and `level`. `n` is the number of observations, and `p` is their dimensionality, and generate `X` as an $n \times p$ matrix of mean-0, variance-1 Gaussian elements. The weight vector w is a p -dimensional vector, all of whose elements are 0 except the first `sparsity` elements, which all take value `level`. Generate the output vector `Y` as

$$Y_i = X_i w + \epsilon_i$$

where X_i is the i th input, and ϵ_i is Gaussian noise. Do not use for loops. [10]

- (b) Write a function `lasso_loss` that takes two inputs `w` and `lambda` and returns the values of the LASSO loss function for (X, Y) . You can treat (X, Y) as additional inputs, or as global variables. [5]
- (c) Generate a dataset with `n=50`, `p = 100`, `sparsity=5`, `level=5`. [5]
- (d) Use the `optim` function to find the best-values of w for the dataset above on the LASSO loss function. Set `lambda=1`. Plot the true w and the returned w . [10]
- (e) Use the `optim` function to find the best-values of w and `lambda` for the dataset above on the LASSO loss function. Plot the true w and the returned w . [10]

Now we are going to directly solve the LASSO problem.

2 Problem 2: Coordinate descent [60]

1. First we'll solve the 1-d case. Write a function `lasso1d` that takes three inputs, length- n inputs `x`, `y` and `lambda`, and returns a scalar weight `w` by first calculating the OLS solution (correlation coefficient) and then soft-thresholding it. See the slides. [10]
2. Given a p -dimensional weight vector, write a function `get_residual` to calculate the residual for some dimension `dim`. This function should take two inputs `w` and `dim` (and `X, Y` unless they are global), and return the residual error from trying to predict `Y` using all dimensions of `X` except `dim`. The simplest way to do this is to set `w[dim] <- 0`, and then calculate `Y_pred = X · w`. The residual is the difference between the true `Y` and `Y_pred`. [10]

3. Now we will solve for the p -dimensions \mathbf{w} vector by coordinate descent. Initialize \mathbf{w} to some value. Cycle through each dimension, first calculating its residual, and then updating the corresponding component of \mathbf{w} . Repeat this until the change in \mathbf{w} after an entire sweep is less than some threshold. [20]
4. Try this on your earlier dataset, again with `lambda = 1`. Comment on your solution obtained this way versus the solution obtained from `optim` [10]
5. Rerun your algorithm from the first n elements of X , where n varies from 0 to 50 in steps of 5. Plot the L_2 error between the resulting w and the true w . [10]